
Goldman Sachs × Anthropic

A Strategic Account Brief

Mario Enrique Mercado

Internal-facing brief for Goldman Sachs territory hand-off · Submitted for Anthropic BDR consideration · May 2026

Goldman Sachs at a Glance

~46,000

employees firmwide

10,000+

GS AI Assistant users; firmwide rollout through 2025

Opus 4.6

Claude agents in production for accounting, compliance, client onboarding (Feb 2026)

Embedded

Anthropic engineers inside Goldman tech teams ~6 months (per CNBC, Feb 2026)

\$1.5B JV

Anthropic + Blackstone + H&F + Goldman launch AI services firm (May 4, 2026)

Peers Running Claude, and the Ones Running OpenAI

- *Goldman's four major US peers run firmwide internal AI assistants at 85-90% workforce coverage.*
- *GS AI Assistant active footprint sits at ~22% (10K+ on a 46K firm) despite firmwide rollout in 2025.*
- *Two peers publicly run Anthropic. One is OpenAI-locked.*

Bank	Internal AI Tool	Vendor(s) Named	Scale + Recency
JPMorgan Chase	LLM Suite	OpenAI + Anthropic	~230-250K employees (~85% on LLM Suite); agentic AI 2026
Morgan Stanley	AI @ Morgan Stanley + Debrief + AskResearchGPT	OpenAI (exclusive wealth partner since 2023)	98% advisor adoption; expanded to trading
Citi	Stylus Workspaces + Citi Assist	Google Gemini + Anthropic Claude	~150-182K employees across 11 countries; 5K-employee agentic pilot 2026
Bank of America	Erica for Employees + AskGPS	Undisclosed (in-house build)	~90% of workforce on Erica; 17K developers on coding tool; \$4B AI/tech 2025

The Thesis

- *Anthropic and Goldman partnership: Opus 4.6 agents in production, embedded Anthropic engineers, and a \$1.5B AI-services JV announced (May 4, 2026).*
- *This brief covers internal expansion under Goldman CIO Marco Argenti, framed around the three capabilities he named as the biggest 2026 AI leaps.*

Long-Document Reasoning

Argenti 2026:

"AI models capable of handling far more context."

Anthropic Fit:

Sonnet 4.6 and Opus 4.7 are a native fit for IB document workloads: M&A diligence, S-1 review, and 10-K analysis.

Agentic Execution

Argenti 2026:

"Will function like a computer's operating system... execute multistep tasks."

Anthropic Fit:

Opus 4.6 agents already run on Goldman workloads in compliance and onboarding. Claude Code is the adjacent surface not yet deployed, bringing codebase-aware agentic development to the engineering org.

Firmwide Knowledge Access

Argenti 2026:

"Re-tune and re-tool itself for AI."

Anthropic Fit:

Claude Enterprise extends projects, governance controls, and a no-training guarantee on customer data to all 46,000 employees.

Sizing the Prize

- *Public pricing × Goldman seat populations × Tier 1 enterprise discount levels = \$8-12M ARR mid-case at Year 2 steady state on the two unowned surfaces.*
- *\$16M+ upside if pricing holds near list. Directional only; contract path and final discount are AE-owned.*

Pricing Scenario	Code ARR	Enterprise ARR	Combined
List (\$1,600 / \$800 per seat)	~\$12.9M	~\$3.6M	~\$16.5M
Tier 1 typical (30% off)	~\$9.0M	~\$2.5M	~\$11.6M
Tier 1 aggressive (50% off)	~\$6.4M	~\$1.8M	~\$8.3M

Use Case 1: Claude Code in Goldman Engineering

- *Opus 4.6 agents already run on Goldman's process workloads.*
- *Claude Code is the next surface: codebase-aware, agentic development in the engineer's daily flow.*

“Engineers are no longer primarily in the business of writing code but are increasingly in the business of supervising the machines that write it for them.”

— Marco Argenti, CIO, Goldman Sachs

- Claude Code lives in the engineer's daily flow: codebase-aware, terminal-native, and built into the development loop.
- Model Context Protocol (MCP) connects Claude Code to internal Goldman systems and data, keeping code inside Goldman's trust boundaries.
- ~11,500 engineers firmwide (~25% of 46,000) sit in the target seat population. Pilot scope is a single platform team; expansion economics scale across the engineering org.
- Pilot scoping, success-metric definition, and pricing are AE/SE/FDE-owned. The BDR work product on this surface is candidate workload identification only.

Use Case 2: Claude Enterprise for Investment Banking

Long-context analysis on the workloads that define the analyst job.

- The target workloads (S-1 review, M&A diligence room synthesis, comparable company analysis, earnings prep) are document-heavy and anchor the analyst job.
- IB analysts work 70-100 hours/week, and M&A data rooms routinely contain thousands of documents per deal. Even modest cycle-time compression on synthesis is material.
- Claude Enterprise project context is built for this exact workload shape, handling "ten 30-minute sales transcripts, 15 full financial reports, and 100K lines of code" per workspace.
- Adds firmwide governance features that raw API access cannot provide: SSO/SCIM, audit logs, no-training guarantees on customer data, and knowledge connectors.
- Quantum reports a 90% reduction in business proposal time using Claude Enterprise. Anthropic-published; Goldman MRM should expect to validate independently.

The Vendor Landscape at Goldman

- *Goldman's AI architecture is multi-vendor by design, so the expansion play is workload-by-workload, not vendor displacement.*
- *The competitive surprise on agentic coding is Cognition Devin, not OpenAI.*

Vendor	Footprint at Goldman	Implication for Anthropic
OpenAI / Microsoft Azure	Available in GS AI Assistant: GPT-4o, GPT-4o-mini, o3-mini. GitHub Copilot firmwide for ~12K developers.	Coexist: Anthropic stays in Goldman's multi-model picker. Compete on workload fit: Claude leads on long-context analysis and agentic process workloads.
Google Gemini / Vertex	Available in GS AI Assistant: Gemini 2.0 Flash, 1.5 Pro. Vertex agent platform; GS named lead customer at Google Cloud Next 2026.	Defend: Vertex is becoming Goldman's default agent runtime. Position Claude Code and Claude Enterprise as the agentic surface, not just another model in the picker.
Cognition Devin	Engineering pilot scaling to 'thousands' alongside ~12K developers (per Argenti, July 2025).	Compete: Devin occupies the agentic-coder narrative Argenti has publicly championed. Position Claude Code as the head-to-head alternative.
Anthropic Claude	Available in GS AI Assistant. Opus 4.6 agents in compliance, accounting, onboarding. Embedded engineers. \$1.5B JV May 4, 2026.	Expand: build on Anthropic's existing footprint at Goldman by extending Claude Code into engineering and Claude Enterprise into IB.

Where Claude Wins, Workload by Workload

- *Anthropic already wins on agentic process workloads with Opus 4.6 in production.*
- *The table below maps adjacent surfaces where Anthropic should expand next.*

Workload	Best Fit	Why
Long-context document analysis (M&A, S-1, 10-K)	Claude Sonnet 4.6 / Opus 4.7	Native long-context handling for dense financial documents.
Codebase-aware agentic development	Claude Code	Anthropic-built agentic development tool with MCP access to Goldman's internal systems.
Firmwide governance + projects (non-engineer)	Claude Enterprise	Firmwide governance: SSO/SCIM, audit logs, project context, and a contractual guarantee that Goldman data is not used to train Anthropic models.
General-purpose chat / utility	GS AI Assistant model picker (no change)	Already deployed firmwide; no change needed for general utility.

Objections & Risks

Five objections a senior buyer may raise on a strategic-account expansion play.

Objection	Response
"We already use ChatGPT Enterprise / GPT-4o."	Goldman's AI architecture is multi-model by design, so we're not asking you to switch. We're asking you to expand Claude on the workloads where it wins on the 2026 priorities Argenti has publicly identified.
"What about model risk management and vendor security review?"	Anthropic offers contractual no-training guarantees on customer data, SOC 2 compliance, HIPAA-ready deployments, and availability via AWS Bedrock or Google Cloud Vertex to fit Goldman's existing cloud security posture. Anthropic provides model risk management (MRM) documentation through its enterprise engagement teams.
"How does this get deployed without overwhelming our platform team?"	Anthropic engineers are already embedded inside Goldman tech teams. Expansion extends an existing engagement footprint, not a net-new one. Pilot scoping is AE/SE/FDE-owned; the BDR work product on this surface is candidate workload identification only.
"What's the procurement and security-review timeline for adding a new vendor surface?"	Tier 1 bank vendor onboarding typically runs 6-9 months. Anthropic's existing FDE engagement and SOC 2 / HIPAA-ready posture compress portions of that timeline. The brief assumes a multi-quarter cycle from first meeting to firmwide deployment.
"Claude is in GS AI Assistant's model picker — isn't that already enough?"	Availability is not adoption. If the model picker defaults to non-Claude, Claude usage stays low. Claude Code and Claude Enterprise put Claude on the workload, not behind a router.

Buying Committee Map

Who advances the deal, who stops it. The BDR opens the conversation then the AE takes the meeting and owns the deal cycle.

EXECUTIVE SPONSOR

Marco Argenti, CIO

Functional Economic Buyer for AI direction; firmwide signing routes through CFO / Management Committee. Owns AI strategy publicly and aligns with Anthropic's 2026 thesis.

STRATEGY LEAD

Office of Applied Innovation

George Lee and Jared Cohen co-lead the Office. Public mandate to source AI use cases firmwide; both report directly to CEO David Solomon.

TECHNICAL LEADS

Platform / IB Tech leadership

Engineering leader for Claude Code; IB Tech lead for Claude Enterprise pilot.

GATEKEEPERS

Vendor risk, MRM, infosec, compliance

Owns contract review, model risk documentation, and data handling review.

The Opening Play

Day-one multi-thread communication: one detailed outbound to Cohen, plus two threads to engineering and IB tech leadership.

OUTBOUND MESSAGE — JARED COHEN, CO-HEAD, OFFICE OF APPLIED INNOVATION, GOLDMAN SACHS

Subject: Argenti's ISDA bootcamps and your use-case sourcing

Hi Jared,

Marco Argenti's remarks at ISDA Boston last Thursday on sending nearly half the firm through AI bootcamps stood out for how concretely they framed the workforce side of this.

I'm in Business Development at Anthropic, reaching out on a narrower angle. Anthropic engineers and Goldman tech teams already run Opus 4.6 agents in production for compliance, accounting, and onboarding. The Office of Applied Innovation's mandate to source firmwide AI use cases overlaps directly with how our team partners with strategic accounts, and I wanted to ask where your team is sourcing the next wave.

Hoping to set up 20 minutes on your calendar to walk through a few workload examples from peer institutions. Happy to send them over ahead of time if useful.

PARALLEL THREADS TO CONSIDER

Senior Eng Leader, Argenti's org (TBD on AE intake)

Devin Pilot: Claude Code as the Engineer's Surface

- Hook: Argenti's Jul 2025 Devin announcement scaling to 'thousands' alongside ~12K devs
- Frame: Claude Code is codebase-aware, MCP-integrated, terminal-native; complementary to Devin
- Ask: 30-min scoping with one platform team's modernization workload

IB Technology Leader (TBD on AE intake)

IB Doc Workloads: Claude Enterprise for Long-Context

- Hook: 70-100hr analyst weeks; thousands of docs per M&A data room
- Frame: Claude Enterprise long-context plus governance (SSO, audit, no-train) the picker can't match
- Ask: 30-min scoping with one analyst workflow

NEXT STEP · Book 60-min scoping with AE, Solutions Engineer, and embedded FDE lead at Goldman. Share workload candidates with AE in advance.

First 90 Days on the Goldman Territory

30/60/90: research-led outbound, multi-threaded outreach, and coordination with the AE and FDE team already on the account.

Days 1–30 • Discover & Open

1. Map the AE's active workstreams and the embedded FDE team's footprint at Goldman.
2. Set up signal monitoring on Argenti, GS earnings, public AI announcements, and M&A pipeline news.
3. Launch multi-threaded outbound to Argenti's org and the Office of Applied Innovation, anchored to recent public signals.
4. Target: 1-2 qualified meetings handed to the AE.

Days 31–60 • Multi-Thread & Hand Off

1. Hand off the first qualified meetings to the AE; stay tagged in for follow-up when useful.
2. Expand outbound to IB Tech leadership and platform engineering leads, coordinating with the embedded FDE team.
3. Continue responding to public signal events with tailored outreach.
4. Target: 3-5 multi-threaded conversations active across personas.

Days 61–90 • Sustain & Document

1. Support the AE through scoping; pull in SE, Applied AI, and FDE intros when needed.
2. Qualify a second persona group (research-side AI lead, business-unit tech leads).
3. Document the Goldman territory playbook for future reps and refine the messaging archive.
4. Target: named pipeline coverage at multiple stages on the account, no overlap with AE-active workstreams.

What Needs to Be Confirmed in Week One

Three open questions for the account team in week one. Each affects how the deal runs or which personas to prioritize. Pricing, FDE assignment, and pilot scoping are intentionally off this list, since those are account-team decisions, not BDR intake.

Open Question	Why It Matters
Is Goldman's Anthropic access direct, via AWS Bedrock, or via GCP Vertex?	Direct vs. cloud co-sell changes the deal motion, stakeholder set, and realized ACV. Slide 5's sizing math assumed direct.
Are Nachmann's AWM org (May 4 JV) and Argenti's CIO org (internal expansion) coordinated on Claude positioning?	Determines whether the BDR multi-thread can include AWM-side personas or must explicitly avoid them. Uncoordinated outreach across two buying centers creates cross-thread risk.
What are Anthropic's embedded engineers running at Goldman, and who should the BDR target for slide 12's parallel threads?	Active FDE flight tells the BDR what to avoid duplicating. AE-identified entry points unblock slide 12's parallel threads.